



# PREDICTING DIABETES

Sami Chowdhury, Yiannis Pagkalos, Lauren Christiansen, Mei  
Kam Bharadwaj, Dhvani Patel



# Table of contents



**01**

**Executive  
Summary**

---

**02**

**Goals and  
Questions**

---

**03**

**Data Collection  
and Cleanup**

---

**04**

**Approach/Issues  
and Resolution**

---

**05**

**Results/ Conclusion**

---

**06**

**Potential Next Steps**

---





# EXECUTIVE SUMMARY



- Our **project goal** was to determine how key datapoints (BMI, high blood pressure, cholesterol, stroke, heart disease/attack, physical activity level, general health level, physical health level, difficulty walking scale, age, education level, income level) relate to the diagnosis of diabetes in different patients.



- We used this underlying data to create and train machine learning models to easily predict whether a patient would be diagnosed.





# GOALS/ QUESTIONS



Our aim was to identify and utilize the most impactful datapoints (e.g., BMI, blood pressure, cholesterol, lifestyle factors) to improve our model's accuracy and reliability on predicting a diabetes diagnosis.

## Questions:

1. Which dataset is best to train our model (number of features, rows and quality)?
2. Which features are most strongly correlated to diabetes diagnoses, and how do they contribute to the model's predictions?
3. What are the models that will produce the best accuracy scores?
4. How can we finetune the selected models to amplify accuracy scores?





# DATA COLLECTION / CLEAN UP



- Collected at the following Kaggle link by Alex Teboul
- Obtained from the Behavioral Risk Factor Surveillance System (BRFSS), an annual telephone survey that is collected annually by the CDC.
- The features are either questions asked of participants or variables calculated based on their responses. The dataset includes the following:
  - ['Diabetes\_binary', 'HighBP', 'HighChol', 'CholCheck', 'BMI', 'Smoker', 'Stroke', 'HeartDiseaseorAttack', 'PhysActivity', 'Fruits', 'Veggies', 'HvyAlcoholConsump', 'AnyHealthcare', 'NoDocbcCost', 'GenHlth', 'MentHlth', 'PhysHlth', 'DiffWalk', 'Sex', 'Age', 'Education', 'Income']





# APPROACH



- Data was already balanced, cleaned for nulls (Keggle included both balanced and unbalanced versions of the data)
- Required encoding of a few set features (BMI condensed to scientific classifications from Underweight to Obesity 3, and initially also encoded AGE, INCOME, EDUCATION)
- In the end, we only encoded BMI and kept the original groupings of AGE, INCOME, and EDUCATION as this ended up producing better results
- Initially we did not remove features that were not strongly correlated to the results. Eventually removing those improved our model performance.
- Data was scaled since we had many ordinal encoders, all categorical and ordered
- Finally, we run GridSearchCV for all the models to further optimize performance by identifying model parameters that could help

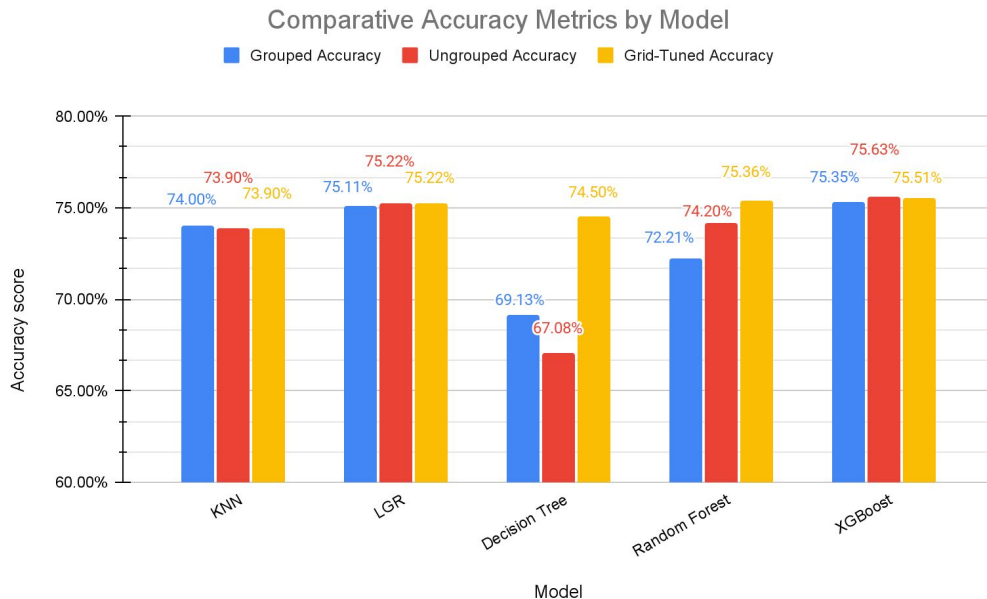


# RESULTS AND CONCLUSIONS

Given the nonlinear classification of the data, we experimented with the following models below. After all optimizations the following results were achieved:

- Decision Tree: 74.5%
- KNeighborsClassifier: 73.9%
- Logistic Regression: 75.2%
- RandomForestClassifier: 75.4
- XGBClassifier: 75.6%

# RESULTS AND CONCLUSIONS



- The biggest improvements were achieved in the Decision Tree and the Random Forest model after the GridSearch optimization
- Interestingly, KNN and LGR models performed slightly worse after we ungrouped the Age, Education and Income categories.
- Similarly, XBoost performed better before GridSearch but only marginally.





# DECISION TREE



	Predicted 0	Predicted 1	<u>Accuracy: 0.67</u>		
Actual 0	6240	2607	Test score: 0.67		
Actual 1	3246	5580	Train Score: 0.93		
Classification Report					
	precision	recall	f1-score	support	
0.0	0.66	0.71	0.68	8847	
1.0	0.68	0.63	0.66	8826	
accuracy			0.67	17673	
macro avg	0.67	0.67	0.67	17673	
weighted avg	0.67	0.67	0.67	17673	

No parameters specified

## GridSearchCV

	Predicted 0	Predicted 1	<u>Accuracy: 0.745</u>		
Actual 0s	6336	2511	Test score: 0.745		
Actual 1s	1995	6831	Train Score: 0.745		
Classification Report					
	precision	recall	f1-score	support	
0.0	0.76	0.72	0.74	8847	
1.0	0.73	0.77	0.75	8826	
accuracy			0.75	17673	
macro avg	0.75	0.75	0.74	17673	
weighted avg	0.75	0.75	0.74	17673	
{ 'max_depth': 7, 'min_samples_leaf': 1, 'min_samples_split': 2 }					
best score 0.738					

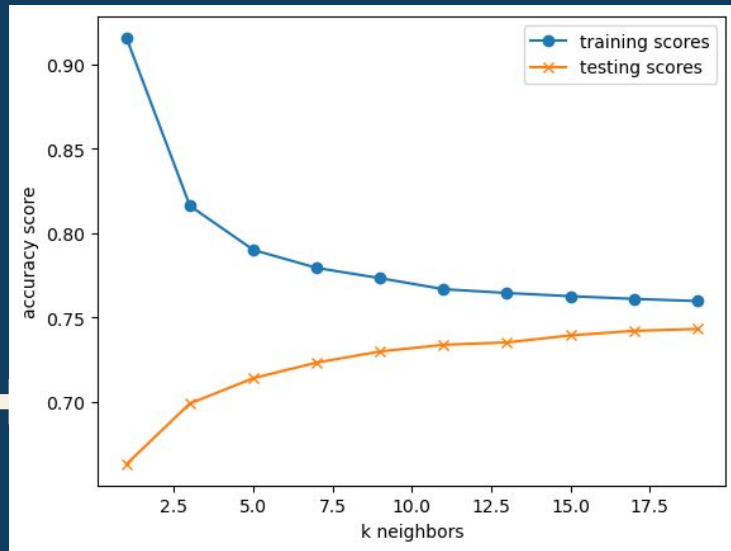


- Specifying the additional parameters (max depth, min sample leaf and min sample split gives a significant improvement in accuracy: - (from 67% to 74.5%)





# KNN MODEL



## Confusion Matrix

	Predicted 0	Predicted 1
Actual 0s	6239	2608
Actual 1s	1997	6829

**Accuracy: 0.739**  
Test score: 0.739  
Train Score: 0.763

## Classification Report

	precision	recall	f1-score	support
0.0	0.76	0.71	0.73	8847
1.0	0.72	0.77	0.75	8826
accuracy			0.74	17673
macro avg	0.74	0.74	0.74	17673
weighted avg	0.74	0.74	0.74	17673



- Training and Test scores converged as k increased
- k=15 seems the best choice for this dataset
- For k=15 accuracy= 73.9%





# KNN MODEL



## GridSearchCV

	Predicted 0	Predicted 1	<u>Accuracy: 0.739</u>	
Actual 0s	6239	2608	Test score: 0.739	
Actual 1s	1997	6829	Train Score: 0.763	
Classification Report				
	precision	recall	f1-score	support
0.0	0.76	0.71	0.73	8847
1.0	0.72	0.77	0.75	8826
accuracy			0.74	17673
macro avg	0.74	0.74	0.74	17673
weighted avg	0.74	0.74	0.74	17673
{ 'algorithm': 'auto', 'metric': 'euclidean', 'n_neighbors': 15, 'weights': 'uniform' }				
best score 0.733				

Additional parameter setting did not improve the model score  
(algorithm, metric, weights)

Max achieved accuracy - 73.9%





# LOGISTIC REGRESSION



## Confusion Matrix

	Predicted 0	Predicted 1
Actual 0s	6513	2334
Actual 1s	2045	6781

Accuracy: 0.752  
Test score: 0.752  
Train Score: 0.745

## Classification Report

	precision	recall	f1-score	support
0.0	0.76	0.74	0.75	8847
1.0	0.74	0.77	0.76	8826
accuracy			0.75	17673
macro avg	0.75	0.75	0.75	17673
weighted avg	0.75	0.75	0.75	17673

```
random_state=1, max_iter=100
```

## GridSearchCV

	Predicted 0	Predicted 1
Actual 0s	6513	2334
Actual 1s	2045	6781

Accuracy: 0.752  
Test score: 0.752  
Train Score: 0.745

## Classification Report

	precision	recall	f1-score	support
0.0	0.76	0.74	0.75	8847
1.0	0.74	0.77	0.76	8826
accuracy			0.75	17673
macro avg	0.75	0.75	0.75	17673
weighted avg	0.75	0.75	0.75	17673

```
{'C': 0.1, 'max_iter': 100, 'penalty': 'l2', 'solver': 'liblinear'}  
best score 0.744
```



- Additional parameter setting did not improve performance:
- Max Accuracy = 75.2%





# RANDOM FOREST



## GridSearchCV

	Predicted 0	Predicted 1	Accuracy: 0.742
Actual 0s	6393	2454	Test score: 0.742
Actual 1s	2112	6714	Train Score: 0.735

### Classification Report

	precision	recall	f1-score	support
0.0	0.75	0.72	0.74	8847
1.0	0.73	0.76	0.75	8826
accuracy			0.74	17673
macro avg	0.74	0.74	0.74	17673
weighted avg	0.74	0.74	0.74	17673

```
random_state=1, n_estimators=500, max_depth=3
```

	Predicted 0	Predicted 1	Accuracy: 0.754
Actual 0s	6351	2496	Test score: 0.754
Actual 1s	1855	6971	Train Score: 0.776

### Classification Report

	precision	recall	f1-score	support
0.0	0.77	0.72	0.74	8847
1.0	0.74	0.79	0.76	8826
accuracy			0.75	17673
macro avg	0.76	0.75	0.75	17673
weighted avg	0.76	0.75	0.75	17673

```
{'max_depth': 11, 'n_estimators': 1000}
```

- Increasing max depth to 11 from 3 and n\_estimators from 500 to 1000, resulted in a modest increase in accuracy (from 74.2% to 75.4%)





# XGBOOST



	Predicted 0	Predicted 1
Actual 0s	6322	2525
Actual 1s	1782	7044

Accuracy: 0.756  
Test score: 0.756  
Train Score: 0.754

## Classification Report

	precision	recall	f1-score	support
0.0	0.78	0.71	0.75	8847
1.0	0.74	0.80	0.77	8826
accuracy			0.76	17673
macro avg	0.76	0.76	0.76	17673
weighted avg	0.76	0.76	0.76	17673

```
xgb_model = XGBClassifier(random_state=1,  
learning_rate=0.05, n_estimators=1000, max_depth=3)
```

## GridSearchCV

	Predicted 0	Predicted 1
Actual 0s	6317	2530
Actual 1s	1798	7028

Accuracy: 0.755  
Test score: 0.755  
Train Score: 0.752

## Classification Report

	precision	recall	f1-score	support
0.0	0.78	0.71	0.74	8847
1.0	0.74	0.80	0.76	8826
accuracy			0.76	17673
macro avg	0.76	0.76	0.75	17673
weighted avg	0.76	0.76	0.75	17673

```
{'learning_rate': 0.2, 'max_depth': 3, 'n_estimators': 100}  
best score 0.748
```



- Interestingly, GridSearchCV results in a slightly less accurate model. Our initial parameter setting (learning rate = 0.05 vs 2 n\_estimators = 1000 vs 100 and max depth 3 (same) give the best result from all our efforts: Accuracy of 75.6%





# POTENTIAL NEXT STEPS



- We could explore more datasets that include health indicators such as HbA1C (hemoglobin A1C) and fast blood sugar test (FBS).
  - HbA1C test measures the average blood sugar (glucose) level over the past 60-90 days. A fasting blood sugar test measures the blood sugar levels first thing in the morning before the patient breaks their fast.
  - If the patient's blood sugar is high, then it indicates that patient has difficulties breaking down sugar in their body.
  - It is best to look at a dataset that includes both HbA1C and FBS data. HbA1C tests are less sensitive compared to the FBS test, but provides a more comprehensive story on the patient's blood sugar over a period of months.
  - In practice, **both** tests are used in the office to get a more accurate diagnosis of diabetes.
- We could experiment with more models and/or model parameter tuning





**THANK YOU!**







# APPENDIX





# DATA CORRELATION MATRIX

Diabetes_binary	
Diabetes_binary	1.000000
HighBP	0.381516
HighChol	0.289213
CholCheck	0.115382
BMI	0.293373
Smoker	0.085999
Stroke	0.125427
HeartDiseaseorAttack	0.211523
PhysActivity	-0.158666
Fruits	-0.054077
Veggies	-0.079293
HvyAlcoholConsump	-0.094853
AnyHealthcare	0.023191
NoDocbcCost	0.040977
GenHlth	0.407612
MentHlth	0.087029
PhysHlth	0.213081
DiffWalk	0.272646
Sex	0.044413
Age	0.278738
Education	-0.170481
Income	-0.224449





# COLUMNS / COLUMN DESCRIPTIONS



## Columns

In [24]: `diabetes.columns`

Out[24]: Index(['Diabetes\_binary', 'HighBP', 'HighChol', 'CholCheck', 'BMI', 'Smoker',  
'Stroke', 'HeartDiseaseorAttack', 'PhysActivity', 'Fruits', 'Veggies',  
'HvyAlcoholConsump', 'AnyHealthcare', 'NoDocbcCost', 'GenHlth',  
'MentHlth', 'PhysHlth', 'DiffWalk', 'Sex', 'Age', 'Education',  
'Income'],  
dtype='object')

## Column Descriptions

1. Diabetes\_binary: 0 = no diabetes 1 = prediabetes or diabetes
2. HighBP: 0 = no high BP 1 = high BP
3. HighChol: 0 = no high cholesterol 1 = high cholesterol
4. CholCheck: 0 = no cholesterol check in 5 years 1 = yes cholesterol check in 5 years
5. BMI: Body Mass Index





# BMI CLASSIFICATION



## BMI

```
def BMI_classification(x):  
    if x < 18.5:  
        return "Underweight"  
    elif x > 18.5 and x <= 24.9:  
        return "Normal"  
    elif x > 24.9 and x <= 29.9:  
        return "Overweight"  
    elif x > 29.9 and x <= 34.9:  
        return "Obesity 1"  
    elif x > 34.9 and x <= 39.9:  
        return "Obesity 2"  
    elif x > 39.9:  
        return "Obesity 3"
```

```
diabetes_df['BMI'] = diabetes_df['BMI'].apply(BMI_classification)  
diabetes_df['BMI'].value_counts()
```

```
BMI  
Overweight      24135  
Obesity 1       17301  
Normal          14460  
Obesity 2        8112  
Obesity 3        6031  
Underweight      653  
Name: count, dtype: int64
```

